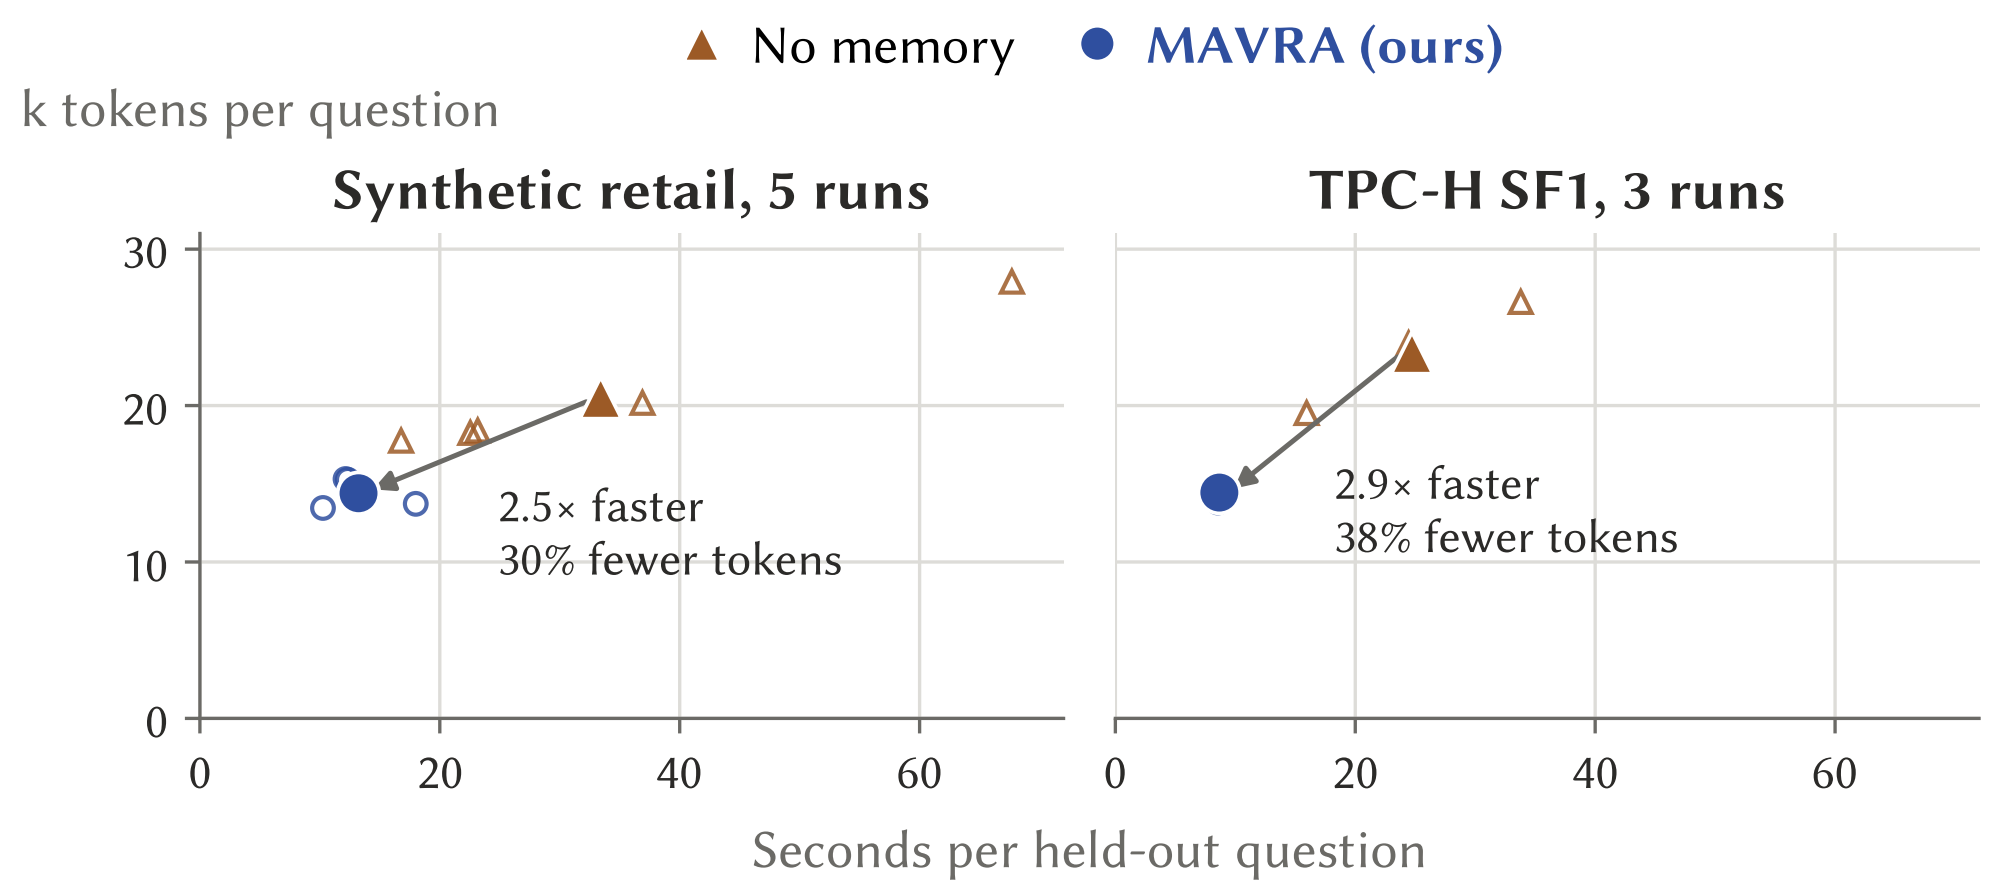
No memory
MAVRA (ours)
k tokens per question
Synthetic retail, 5 runs
TPC-H SF1, 3 runs
30
20
10
0
0
20
40
60
2.5× faster
30% fewer tokens
2.9× faster
38% fewer tokens
Seconds per held-out question